

# AI Agent – Actionable Build Plan (Two-Repo Model)

This document now assumes **Option A**: the existing company-website repo on **Vercel** stays untouched, while a **new backend repo** (**accessibility-scanner-backend**) is created to host the API service, background worker, queue and database—all deployed to **Railway**. The steps below describe only the **backend repo**. A short integration guide for calling the API from the website is included at the end.\*

## Phase 0 – Meta Setup (½ day)

#	Task	Output	Acceptance Criteria
0.1	Create GitHub repo <b>accessibility-scanner-backend</b> (public/private).	Repo initialised	MIT LICENSE & <b>README.md</b> scaffolded
0.2	Enable Issues & project board (Backlog → In Progress → Done).	Project board	Board seeded with this plan's tasks
0.3	Configure Husky with lint-staged & Conventional Commits.	<b>.husky/</b> scripts	<b>git commit</b> blocked on lint failure

## Phase 1 – Repo Scaffolding (1day)

#	Task	Directory	Key Files	Acceptance Criteria
1.1	Init <b>pnpm workspace</b>	/	<b>package.json</b> , <b>pnpm-workspace.yaml</b>	<b>pnpm install</b> succeeds
1.2	Create apps: <b>apps/api</b> , <b>apps/worker</b>	folders	baseline <b>package.json</b> , TS config	<b>pnpm -r build</b> passes
1.3	Add root ESLint, Prettier, TypeScript v5 strict	<b>/.eslintrc.js</b>	<b>pnpm lint</b> 0 errors	

## Phase 2 – Infrastructure as Code (1½ days)

#	Task	Files	Acceptance Criteria
2.1	Docker Compose (local): <b>api</b> , <b>worker</b> , <b>redis</b> , <b>postgres</b>	<b>docker-compose.yml</b>	<b>docker compose up</b> healthy

#	Task	Files	Acceptance Criteria
2.2	<b>Railway Terraform</b> module provisioning Service (api), Job (worker), Postgres, Redis.	<code>infra/railway.tf</code>	<code>terraform plan</code> no drift
2.3	<b>GitHub Actions CI</b> – on push:• install deps• type-check & jest• docker build api & worker• push to <b>GHCR</b> • deploy to Railway via <code>railway up</code> .	<code>.github/workflows/ci.yml</code>	PR green checks

## Phase 3 – Backend API (3days)

#	Task	Endpoint	Acceptance Criteria
3.1	Scaffold <b>Express</b> w/ <code>ts-node</code> in <code>apps/api</code> . Health <code>/health</code> .	Returns <code>{ ok: true }</code>	
3.2	Add <b>Zod</b> schemas for request/response.	<code>src/validation/</code>	Invalid URL → 400
3.3	<b>**POST **/api/scan-website</b> → enqueue BullMQ job, return <code>{ jobId }</code> (UUID).	Unit test passes	
3.4	<b>**GET **/api/scan-status/:jobId</b> returns <code>`QUEUED`</code>	RUNNING	COMPLETED FAILED`. Enum validated
3.5	<b>**GET **/api/scan-results/:jobId</b> returns aggregated results from Postgres.	Mock until Phase 4	
3.6	Add <b>rate-limiter</b> 10 req/min/IP using Redis backend.	Exceeding limit → 429	

## Phase 4 – Worker Service (4days)

#	Task	Description	Acceptance Criteria
4.1	Bootstrap <code>apps/worker</code> – connect to BullMQ queue.	Worker logs 'ready'.	
4.2	On job: launch <b>Puppeteer</b> (headless Chromium) with 30s nav timeout.	<code>page.title()</code> non-empty	

#	Task	Description	Acceptance Criteria
4.3	Inject <b>axe-core</b> ; run scan.	Violations array returned	
4.4	Compute <b>score</b> 0-100 via severity weights.	Unit test: 0 violations → 100	
4.5	Persist to Postgres ( <b>jobs</b> , <b>scans</b> , <b>issues</b> ) via Prisma.	Rows exist	
4.6	ACK or mark FAILED; auto-retry ≤2.	Failed job visible in Bull UI	
4.7	Concurrency via env <b>MAX_WORKERS</b> (default 2).	Parallelism scales	

## Phase 5 – Observability (¾ day)

#	Task	Tech	Acceptance Criteria
5.1	<b>Pino</b> structured logs to Railway Log viewer.	JSON lines displayed	
5.2	<b>Prometheus exporter</b> ( <b>/metrics</b> ) exposing job latency, queue depth.	Curl returns metrics	
5.3	<b>Grafana Cloud</b> free tier: scrape/export; dashboard in <b>observability/</b> .	Graph shows <b>scan_duration_seconds</b>	

## Phase 6 – Documentation & DX (¾ day)

#	Task	Output	Acceptance Criteria
6.1	Generate <b>OpenAPI 3</b> spec ( <b>openapi.yaml</b> ).	Swagger UI renders	
6.2	Update <b>README.md</b> with local dev, Railway deploy, env vars.	New dev up in ≤15 min	
6.3	Seed <b>Changelog.md</b> (Keep-a-Changelog).	Version 0.1.0 entry	

## Phase 7 – Deployment & Smoke Test (½ day)

#	Task	Acceptance Criteria
7.1	<b>CI deploy</b> : build & push images → Railway (api & worker).	Services healthy
7.2	Run smoke scan against <b>https://example.com</b> .	Status completes; JSON returned

## Website-Repo Integration Guide (½ day)

Step	File	Action
1	<code>src/routes/+page.svelte</code> (landing)	<code>fetch("https://api.your-scanner.com/api/scan-website", { method: "POST", body: JSON.stringify({ url }) })</code> → receive <code>jobId</code> → navigate to <code>/scanResults/[jobId]</code> .
2	<code>src/routes/scanResults/+page.svelte</code>	On load, poll <code>/api/scan-status/{jobId}</code> every 3s; once <b>COMPLETED</b> , fetch <code>/api/scan-results/{jobId}</code> and render dashboard.
3	Add <code>.env.public</code> key <code>VITE_SCANNER_API=https://api.your-scanner.com</code> .	Build passes; runtime uses env

*No backend code lives in the website repo; only these two fetch calls are added.*

## Additional Non-Functional Notes

- **Hosting:** API & worker on Railway containers; Redis & Postgres Railway add-ons. Frontend remains on Vercel (no changes to its pipeline).
- **Queue:** BullMQ on Redis Streams.
- **Data retention:** 30-day rolling delete cron (Railway).
- **Concurrency limit:** Start 2 workers (~50 scans max). Scale via Railway slider.
- **Budget:** Railway starter + existing Vercel plan; no new cost on website repo.

## Definition of Done

- User visits company site → submits URL → sees real scan results in ≤60s.
- Backend repo: tests & lint pass; automatic deploy to Railway.
- Website repo: env var points to API; simple fetch integration works.