

# Digital Accessibility Centre - Accessibility Scanner Backend

---

This repository contains the backend services for the DAC's web accessibility scanner application. It's a Node.js-based monorepo that provides an API for initiating scans, a worker for processing them, and a database for storing the results.

## Table of Contents

- [Architecture](#)
  - [Request Flow](#)
- [Technology Stack](#)
- [Project Structure](#)
- [Setup and Installation](#)
  - [Prerequisites](#)
  - [Running Locally](#)
  - [Environment Variables](#)
- [Database Migrations](#)
- [API Endpoints](#)
- [Key Logic](#)
  - [Scan Processing](#)
  - [Accessibility Scoring](#)

## Architecture

The backend is designed as a set of containerized microservices that communicate asynchronously, making it scalable and resilient.

- **API Service ([apps/api](#)):** A Node.js/Express.js application that serves as the public-facing entry point. It receives scan requests from the frontend, adds them to a job queue, and provides endpoints to retrieve scan results.
- **Worker Service ([apps/worker](#)):** A Node.js application that listens for jobs on the queue. It performs the heavy lifting: launching a headless browser (Puppeteer) to visit the target URL, running [axe-core](#) to perform the accessibility audit, taking a screenshot, and saving the results to the database.
- **Redis / BullMQ:** A Redis-backed message queue (BullMQ) that decouples the API from the Worker. This prevents API requests from timing out while waiting for a long-running scan and allows for job retries and queuing.
- **PostgreSQL Database:** A relational database used to persist all scan results. The schema is managed via migrations.
- **Docker:** All services are containerized with Docker, and [docker-compose.yml](#) orchestrates the local development environment.

## Request Flow

1. A user submits a URL to the frontend (SvelteKit app).
2. The frontend sends a **POST** request to the **API Service** (`/api/scan-website`).
3. The **API Service** generates a unique `jobId` and enqueues a new job in **BullMQ**. It immediately returns the `jobId` to the frontend.
4. The frontend begins polling the **API Service** using the `jobId` (`/api/scan-results/:jobId`).
5. The **Worker Service**, listening to the queue, picks up the job.
6. The **Worker** uses Puppeteer and `axe-core` to scan the site, takes a screenshot, and saves the results (or any errors) to the **PostgreSQL** database, referencing the `jobId`.
7. On a subsequent poll, the **API Service** finds the completed job results in the database and returns them to the frontend.
8. The frontend displays the results to the user.

## Technology Stack

- **Backend:** Node.js, TypeScript, Express.js
- **Database:** PostgreSQL
- **Queue:** Redis, BullMQ
- **Browser Automation:** Puppeteer
- **Accessibility Engine:** `axe-core`
- **Containerization:** Docker, Docker Compose
- **Package Manager:** pnpm (in a monorepo setup)
- **Migrations:** `node-pg-migrate`

## Project Structure

This is a `pnpm` monorepo. The main components are:

```
.
├── apps/
│   ├── api/           # Express API service
│   └── worker/        # BullMQ worker service
├── migrations/        # node-pg-migrate database migration files
├── docker-compose.yml # Local development environment setup
├── package.json       # Root package.json
└── pnpm-workspace.yaml # Defines the pnpm workspace
```

## Setup and Installation

### Prerequisites

- [Docker](#)
- [Docker Compose](#)
- [Node.js](#) (v18 or later)
- `pnpm` (`npm install -g pnpm`)

### Running Locally

1. **Clone the repository:**

```
git clone <repository-url>
cd accessibility-scanner-backend
```

2. **Install dependencies:** This command installs dependencies for all workspace packages (**api** and **worker**).

```
pnpm install
```

3. **Create environment file:** Copy the example environment file and fill in any necessary values (the defaults in **docker-compose.yml** are usually sufficient for local development).

```
# This project does not currently use a .env file, relying on docker-
compose.yml for environment variables.
# If a .env file is added, it should be created from a .env.example.
```

4. **Start the services:** This command will build the Docker images (if they don't exist) and start all services (**api**, **worker**, **postgres**, **redis**).

```
docker compose up --build
```

The API will be available at <http://localhost:3000>.

## Environment Variables

Environment variables are defined in **docker-compose.yml** for local development. Key variables include:

- **POSTGRES\_USER**, **POSTGRES\_PASSWORD**, **POSTGRES\_DB**: Credentials for the PostgreSQL database.
- **POSTGRES\_HOST**, **POSTGRES\_PORT**: Connection details for PostgreSQL (e.g., **postgres**, **5432**).
- **REDIS\_HOST**, **REDIS\_PORT**: Connection details for Redis (e.g., **redis**, **6379**).
- **DATABASE\_URL**: The full connection string for the database, required by **node-pg-migrate**.  
Format: **postgres://<user>:<password>@<host>:<port>/<db>**

## Database Migrations

Database schema changes are managed with **node-pg-migrate**.

- **To run migrations (apply new changes):** The **migrate:up** script is defined in the root **package.json**. It executes the migrations inside the running **api** container.

```
pnpm -w run migrate:up
```

- **To create a new migration:**

```
pnpm -w run migrate:create <migration-name>
```

## API Endpoints

All endpoints are prefixed with `/api`.

---

### POST `/api/scan-website`

Initiates a new accessibility scan for a given URL.

- **Request Body:**

```
{
  "url": "https://example.com"
}
```

- **Success Response (202 Accepted):**

```
{
  "message": "Scan request accepted and enqueued. Job ID: <uuid>",
  "jobId": "<uuid>",
  "submittedUrl": "https://example.com"
}
```

- **Error Response (400 Bad Request):** If the URL is invalid.
- 

### GET `/api/scan-results/:jobId`

Retrieves the results of a scan. This endpoint is designed to be polled by the frontend until the scan is complete.

- **URL Parameters:**
  - `jobId` (string, required): The ID returned from the `POST /api/scan-website` call.
- **Success Response (200 OK):** Returns the full scan result object from the database once the job is complete.

```
{
  "id": 1,
  "job_id": "<uuid>",
  "original_job_id": "<uuid>",
  "submitted_url": "https://example.com",
}
```

```

    "actual_url": "https://example.com/",
    "scan_timestamp": "2023-10-27T10:00:00.000Z",
    "page_title": "Example Domain",
    "scan_success": true,
    "violations": [
        // Array of axe-core violation objects
    ],
    "error_message": null,
    "page_screenshot": "data:image/jpeg;base64,..."
}

```

- **Pending Response (404 Not Found):** If the job is still processing or does not exist, a 404 is returned, and the client should continue polling.

---

## GET /api/export-report/:jobId

Generates and returns a full HTML report for a completed scan.

- **URL Parameters:**
  - **jobId** (string, required): The ID of a completed scan.
- **Success Response (200 OK):**
  - **Content-Type:** `text/html`
  - The response body is a full HTML document containing the formatted scan report, including summary metrics, a screenshot, and a detailed list of all violations.

## Key Logic

### Scan Processing

The worker (`apps/worker/src/index.ts`) contains the core scanning logic:

1. A job is received from the BullMQ queue.
2. Puppeteer launches a headless Chromium browser.
3. The browser navigates to the `submittedUrl`.
4. The `axe-core` script is injected into the page.
5. `axe.run()` is executed to get the accessibility violations.
6. `page.screenshot()` is called to capture a base64-encoded screenshot of the page.
7. The results, including the page title, final URL, violations array, and screenshot, are saved to the `scan_results` table in the PostgreSQL database.
8. If any step fails, the error is caught, and the `error_message` and `scan_success: false` are recorded in the database.

### Accessibility Scoring

The accessibility score is not calculated on the backend but is derived on the **frontend** (`DigitalAccessibilityCentre/src/routes/scan-results/+page.svelte`) for display purposes. The logic is as follows:

- A starting score of 100 is assumed.

- Points are deducted for each violation based on its impact level and the number of nodes it affects:
  - **critical**: 10 points per failing node.
  - **serious**: 5 points per failing node.
  - **moderate**: 2 points per failing node.
  - **minor**: 1 point per failing node.
- The final score is capped between 0 and 100.